

A decorative graphic on the left side of the slide consisting of two overlapping parallelograms. The front one is blue and the back one is a light green. They are positioned diagonally, with the blue one partially covering the green one.

Tensorflow.js


Machine Learning in Javascript

June 6, 2018
Kevin Scott

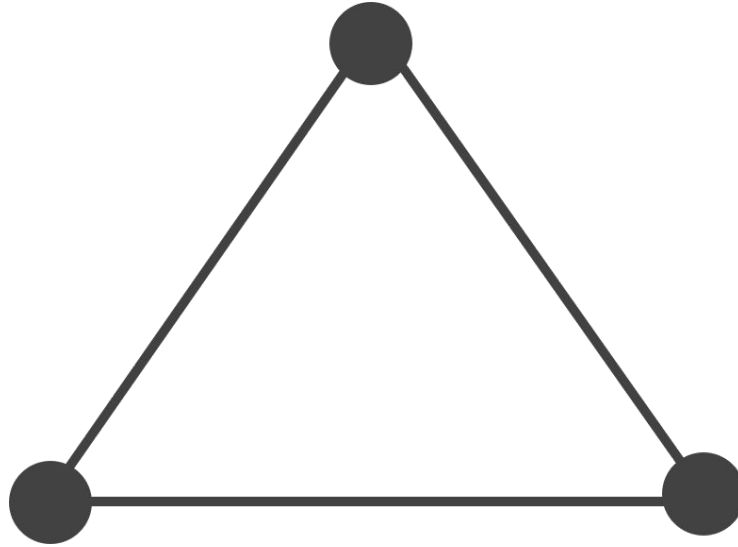
Reboot and Select proper Boot device
or Insert Boot Media in selected Boot device and press a key



Machine Learning

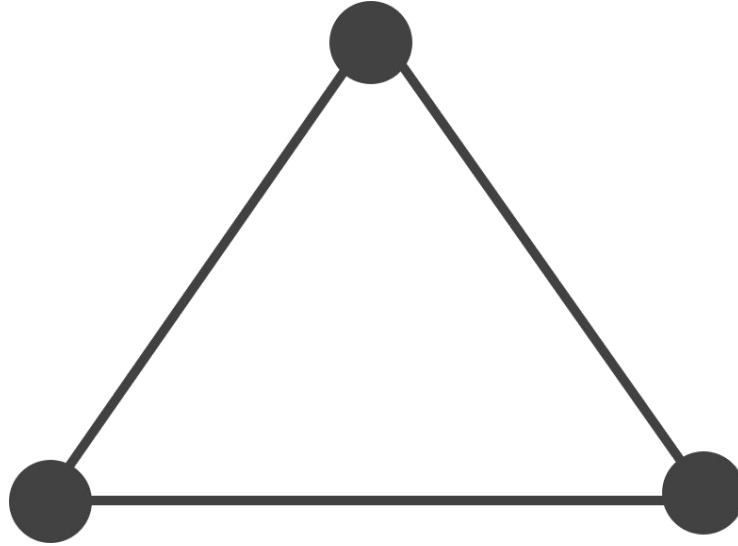


Machine Learning



Mathematics,
Statistics

Machine Learning



Mathematics,
Statistics

Python, R





Why run Tensorflow in your browser?

1. Wide distribution
2. Interactive
3. Sensors
4. Privacy

API

Keras Model

Layers API

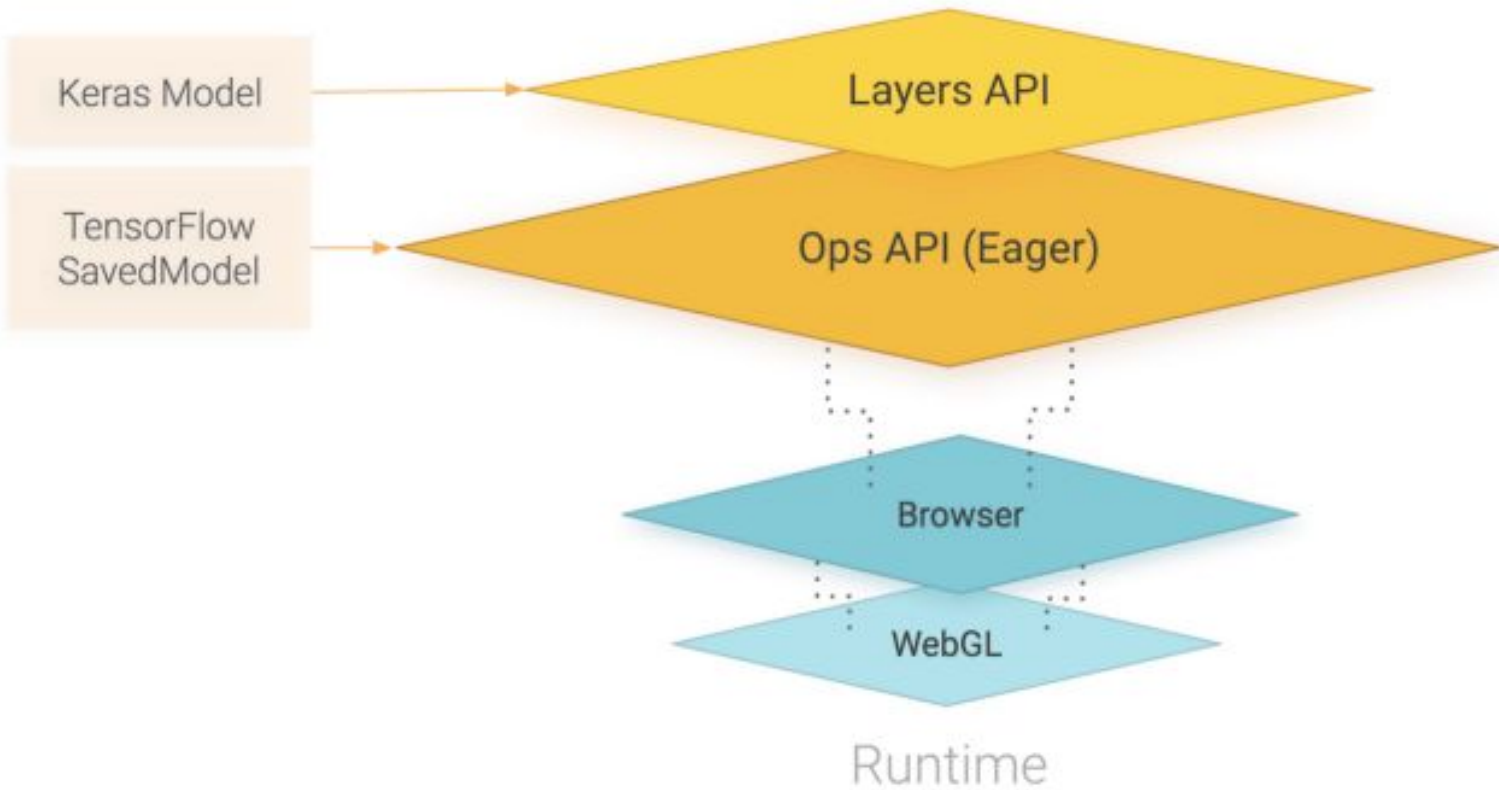
TensorFlow
SavedModel

Ops API (Eager)

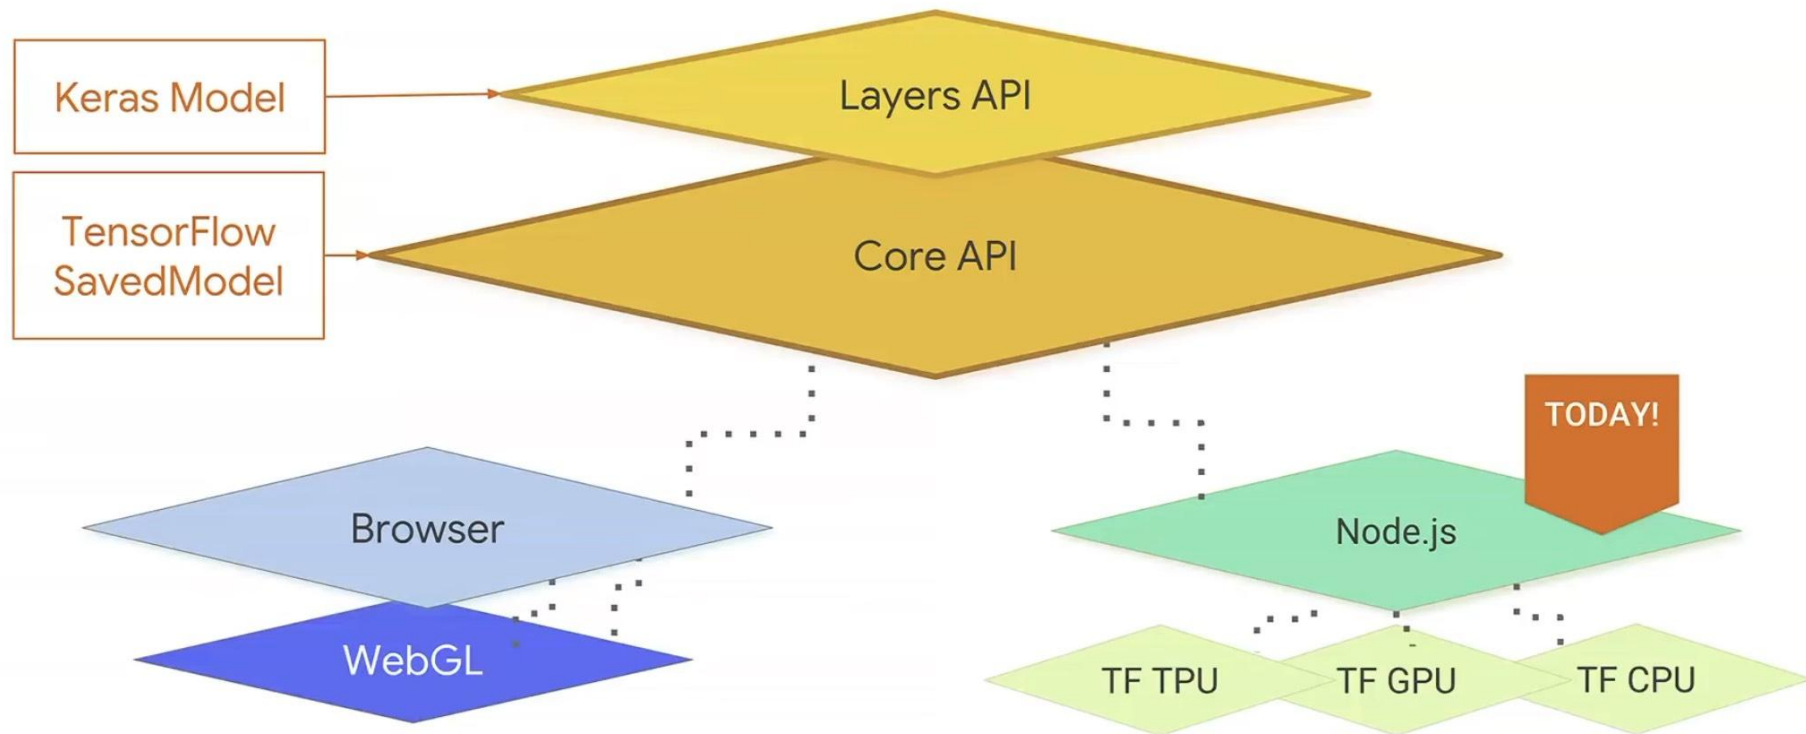
Browser

WebGL

Runtime




API



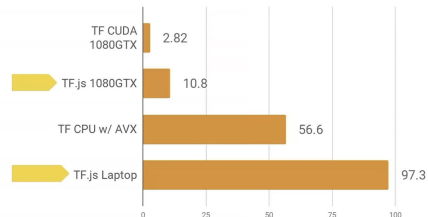
Runtime



Node.js Benchmarks



Inference Time (ms) of MobileNet 1.0_224
Average of 200 runs



Inference Time (ms) of MobileNet 1.0_224
Average of 200 runs



Jeremy Chone



Other recipients: nsth...@google.com, nkats...@gmail.com

Thanks.

Yes, did the test during training and at the end and I am getting a very similar result.

tfjs-node:

training step 0, training accuracy: 0.18 (0.1s)
training step 100, training accuracy: 0.92 (2.1s)
training step 200, training accuracy: 0.94 (2.1s)
training step 300, training accuracy: 1 (2s)
training step 400, training accuracy: 0.92 (2s)
training step 500, training accuracy: 0.98 (2s)
training step 600, training accuracy: 0.94 (2s)
training step 700, training accuracy: 0.98 (2s)
training step 800, training accuracy: 0.92 (2s)
training step 900, training accuracy: 0.98 (2s)

Test set accuracy (batch 1000): 97.00%

python:

step 0, training accuracy 0.14 (0s)
step 100, training accuracy 0.88 (8s)
step 200, training accuracy 0.98 (8s)
step 300, training accuracy 0.94 (8s)
step 400, training accuracy 0.92 (8s)
step 500, training accuracy 0.96 (8s)
step 600, training accuracy 0.96 (8s)
step 700, training accuracy 0.96 (8s)
step 800, training accuracy 1 (8s)
step 900, training accuracy 0.96 (8s)

test accuracy 0.9665



Importing Models

You can import pre-trained models

```
import * as tf from '@tensorflow/tfjs';  
  
const model = await tf.loadModel('http://foo.bar/tfjs_artifacts/model.json');  
// Now the model is ready for inference, evaluation or re-training.
```



Epoch
000,000

Learning rate
0.03

Regularization
None

Regularization rate
0

Problem type
Classification

DATA

Which dataset do you want to use?




Ratio of training to test data: 50%



Noise: 0





Batch size: 10



REGENERATE

FEATURES

Which properties do you want to feed in?



+

-

2 HIDDEN LAYERS

+

-

5 neurons



This is the output from one **neuron**.
Hover to see it larger.

+

-

2 neurons

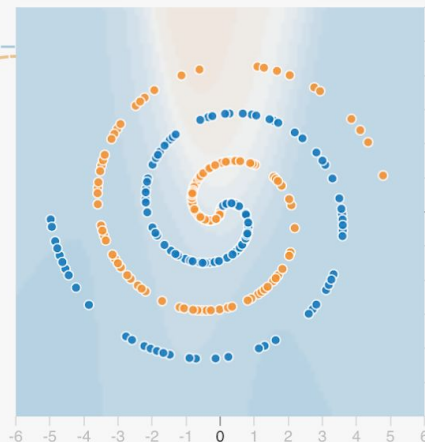


The outputs are mixed with varying **weights**, shown by the thickness of the lines.


OUTPUT

Test loss 0.519

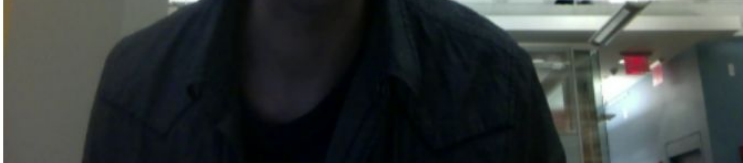
Training loss 0.513



Colors shows data, neuron and weight values.



☐ Discretize output



Content

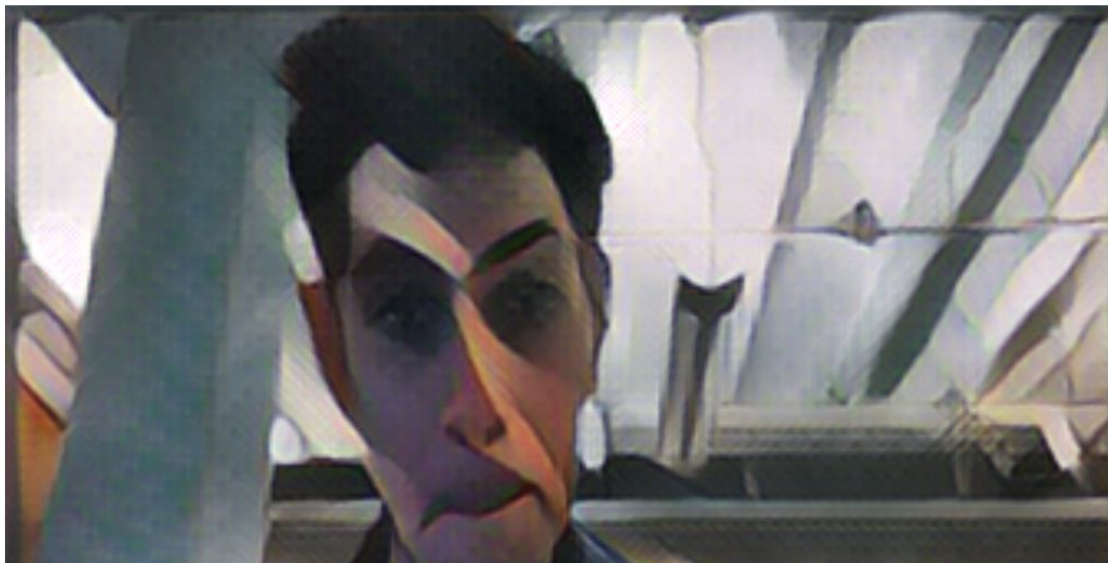
Use webcam ▼



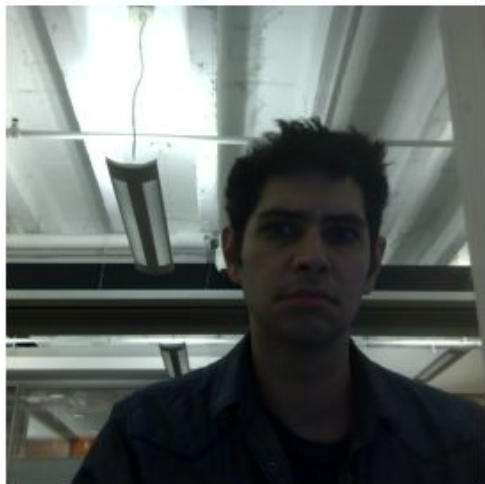
Style

Udnie, Francis Picabia ▼

START STYLE TRANSFER



INPUT



LEARNING

22 EXAMPLES



CONFIDENCE

100%

TRAIN GREEN

21 EXAMPLES



CONFIDENCE

TRAIN PURPLE

32 EXAMPLES



CONFIDENCE

TRAIN ORANGE

OUTPUT

[GIF](#)

Sound

Speech







Thanks!

Google <https://js.tensorflow.org>

<https://github.com/tensorflow/tfjs-examples>

<https://playground.tensorflow.org>

Me <https://thekevinscott.com/slides/tfjs>

<https://twitter.com/thekevinscott>

<https://thekevinscott.com>

Demos [Style transfer in browser](#)

[Teachable Machine](#)

[Pose Estimator](#)